

CSC D70: Compiler Optimization Prefetching

Prof. Gennady Pekhimenko

University of Toronto


Winter 2021

*The content of this lecture is adapted from the lectures of
Todd Mowry and Phillip Gibbons*

No Lecture on April 5th

- I will be involved in MLSys 2021 conference

The Memory Latency Problem



- \uparrow processor speed \gg \uparrow memory speed
- caches are not a panacea

Prefetching for Arrays: Overview

- Tolerating Memory Latency
- Prefetching Compiler Algorithm and Results
- Implications of These Results

Coping with Memory Latency

Reduce Latency:

— Locality Optimizations

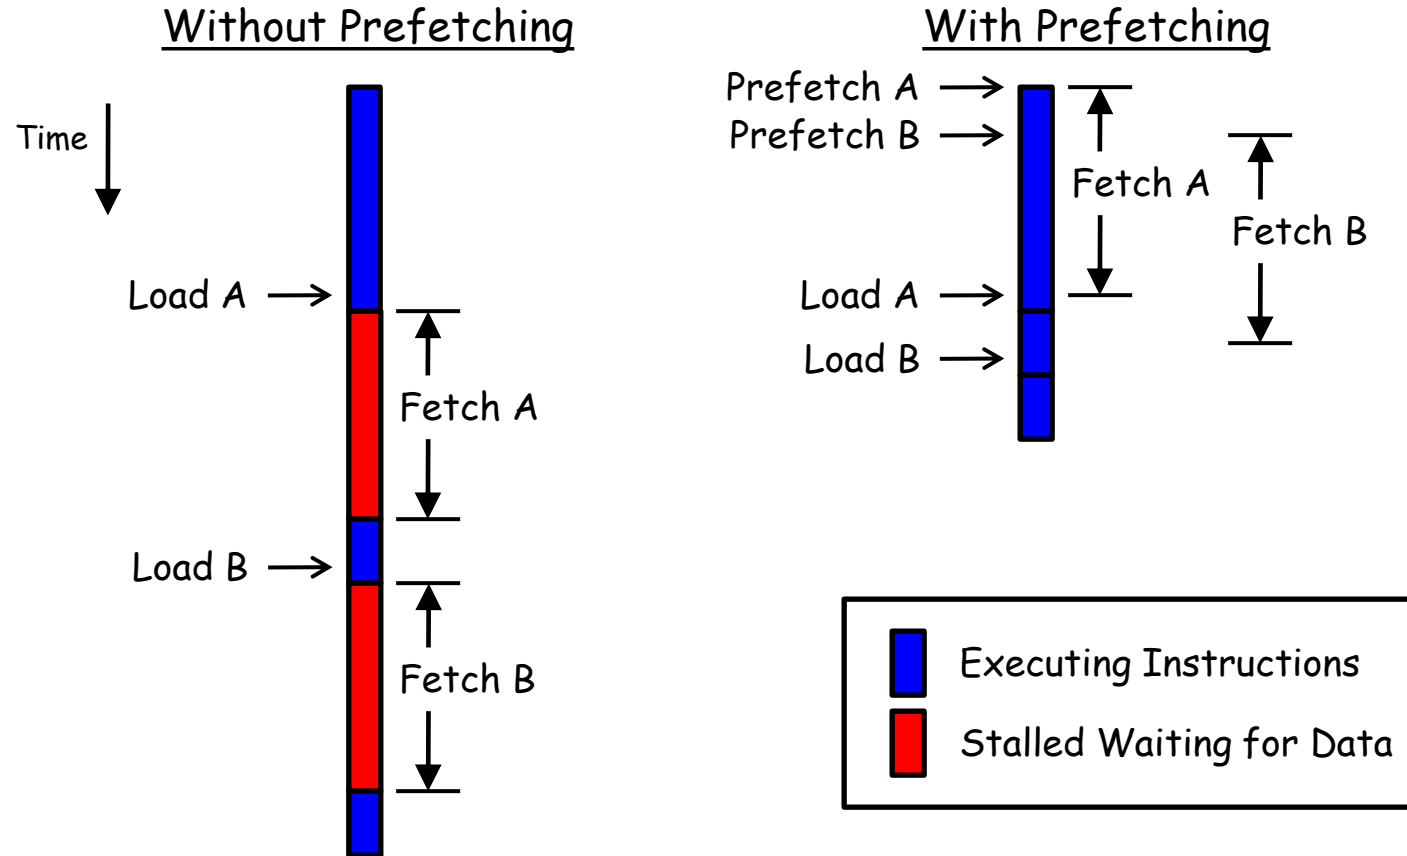
- reorder iterations to improve cache reuse

Tolerate Latency:

— Prefetching

- move data close to the processor before it is needed

Tolerating Latency Through Prefetching



- overlap memory accesses with computation and other accesses

Types of Prefetching

Cache Blocks:

- (-) limited to unit-stride accesses

Nonblocking Loads:

- (-) limited ability to move back before use

Hardware-Controlled Prefetching:

- (-) limited to constant-strides and by branch prediction
- (+) no instruction overhead

Software-Controlled Prefetching:

- (-) software sophistication and overhead
- (+) minimal hardware support and broader coverage

Prefetching Goals

- Domain of Applicability
- Performance Improvement
 - maximize benefit
 - minimize overhead

Prefetching Concepts

possible only if addresses can be determined ahead of time

coverage factor = fraction of misses that are prefetched

unnecessary if data is already in the cache

effective if data is in the cache when later referenced

Analysis: what to prefetch


- maximize coverage factor
- minimize unnecessary prefetches

Scheduling: when/how to schedule prefetches

- maximize effectiveness
- minimize overhead per prefetch

Reducing Prefetching Overhead

- instructions to issue prefetches
- extra demands on memory system



- important to minimize unnecessary prefetches

Compiler Algorithm

Analysis: what to prefetch

- Locality Analysis

Scheduling: when/how to issue prefetches

- Loop Splitting
- Software Pipelining

Steps in Locality Analysis

1. Find data reuse

- if caches were infinitely large, we would be finished

2. Determine “localized iteration space”



- set of inner loops where the data accessed by an iteration is expected to fit within the cache

3. Find data locality:


- $\text{reuse} \cap \text{localized iteration space} \Rightarrow \text{locality}$

Data Locality Example


```
for i = 0 to 2
  for j = 0 to 100
    A[i][j] = B[j][0] + B[j+1][0];
```



Spatial



Temporal



Group

Reuse Analysis: Representation

```
for i = 0 to 2
  for j = 0 to 100
    A[i][j] = B[j][0] + B[j+1][0];
```

- Map n loop indices into d array indices via array indexing function:

$$\vec{f}(\vec{i}) = H\vec{i} + \vec{c}$$

$$A[i][j] = A\left(\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} i \\ j \end{bmatrix} + \begin{bmatrix} 0 \\ 0 \end{bmatrix}\right)$$

$$B[j][0] = B\left(\begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} i \\ j \end{bmatrix} + \begin{bmatrix} 0 \\ 0 \end{bmatrix}\right)$$

$$B[j+1][0] = B\left(\begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} i \\ j \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \end{bmatrix}\right)$$

Finding Temporal Reuse

- Temporal reuse occurs between iterations \vec{i}_1 and \vec{i}_2 whenever:

$$H\vec{i}_1 + \vec{c} = H\vec{i}_2 + \vec{c}$$

$$H(\vec{i}_1 - \vec{i}_2) = \vec{0}$$


- Rather than worrying about individual values \vec{i}_1 of \vec{i}_2 and, we say that reuse occurs along **direction** \vec{r} **vector** when:

$$H(\vec{r}) = \vec{0}$$

- Solution**: compute the *nullspace* of H

Temporal Reuse Example

```
for i = 0 to 2
  for j = 0 to 100
    A[i][j] = B[j][0] + B[j+1][0];
```



- Reuse between iterations (i_1, j_1) and (i_2, j_2) whenever:

$$\begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} i_1 \\ j_1 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} i_2 \\ j_2 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} i_1 - i_2 \\ j_1 - j_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

- True whenever $j_1 = j_2$, and regardless of the difference between i_1 and i_2 .
 - i.e. whenever the difference lies along the nullspace of $\begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}$
 - which is $\text{span}\{(1,0)\}$ (i.e. the outer loop).

Prefetch Predicate

Locality Type	Miss Instance	Predicate
None	Every Iteration	True
Temporal	First Iteration	$i = 0$
Spatial	Every l iterations (l = cache line size)	$(i \bmod l) = 0$

Example: for $i = 0$ to 2
 for $j = 0$ to 100
 $A[i][j] = B[j][0] + B[j+1][0];$

Reference	Locality	Predicate
$A[i][j]$	$\begin{bmatrix} i \\ j \end{bmatrix} = \begin{bmatrix} \text{none} \\ \text{spatial} \end{bmatrix}$	$(j \bmod 2) = 0$
$B[j+1][0]$	$\begin{bmatrix} i \\ j \end{bmatrix} = \begin{bmatrix} \text{temporal} \\ \text{none} \end{bmatrix}$	$i = 0$

Compiler Algorithm

Analysis: what to prefetch

- Locality Analysis

Scheduling: when/how to issue prefetches

- Loop Splitting
- Software Pipelining

Loop Splitting

- Decompose loops to isolate cache miss instances
 - cheaper than inserting IF statements

Locality Type	Predicate	Loop Transformation
None	True	None
Temporal	$i = 0$	Peel loop i
Spatial	$(i \bmod l) = 0$	Unroll loop i by l

- Apply transformations recursively for nested loops
- Suppress transformations when loops become too large
 - avoid code explosion

Software Pipelining

$$\text{Iterations Ahead} = \left\lceil \frac{l}{s} \right\rceil$$

where l = memory latency, s = shortest path through loop body

Original Loop

```
for (i = 0; i < 100; i++)  
    a[i] = 0;
```

Software Pipelined Loop (5 iterations ahead)


```
for (i = 0; i < 5; i++)          /* Prolog */  
    prefetch(&a[i]);  
  
for (i = 0; i < 95; i++) {      /* Steady State */  
    prefetch(&a[i+5]);  
    a[i] = 0;  
}  
  
for (i = 95; i < 100; i++)      /* Epilog */  
    a[i] = 0;
```

Example Revisited

Original Code

```
for (i = 0; i < 3; i++)
  for (j = 0; j < 100; j++)
    A[i][j] = B[j][0] + B[j+1][0];
```

○ Cache Hit
●● Cache Miss



Code with Prefetching

```
prefetch(&A[0][0]);
for (j = 0; j < 6; j += 2) {
  prefetch(&B[j+1][0]);
  prefetch(&B[j+2][0]);
  prefetch(&A[0][j+1]);
}
for (j = 0; j < 94; j += 2) {
  prefetch(&B[j+7][0]);
  prefetch(&B[j+8][0]);
  prefetch(&A[0][j+7]);
  A[0][j] = B[j][0] + B[j+1][0];
  A[0][j+1] = B[j+1][0] + B[j+2][0];
}
for (j = 94; j < 100; j += 2) {
  A[0][j] = B[j][0] + B[j+1][0];
  A[0][j+1] = B[j+1][0] + B[j+2][0];
}

for (i = 1; i < 3; i++) {
  prefetch(&A[i][0]);
  for (j = 0; j < 6; j += 2)
    prefetch(&A[i][j+1]);
  for (j = 0; j < 94; j += 2) {
    prefetch(&A[i][j+7]);
    A[i][j] = B[j][0] + B[j+1][0];
    A[i][j+1] = B[j+1][0] + B[j+2][0];
  }
  for (j = 94; j < 100; j += 2) {
    A[i][j] = B[j][0] + B[j+1][0];
    A[i][j+1] = B[j+1][0] + B[j+2][0];
  }
}
```

Prefetching Indirections

```
for (i = 0; i<100; i++)  
    sum += A[index[i]];
```

Analysis: what to prefetch

- both dense and **indirect** references
- difficult to predict whether indirections hit or miss

Scheduling: when/how to issue prefetches

- **modification of software pipelining algorithm**

Software Pipelining for Indirections

Original Loop

```
for (i = 0; i<100; i++)  
    sum += A[index[i]];
```

Software Pipelined Loop (5 iterations ahead)

```
for (i = 0; i<5; i++)      /* Prolog 1 */  
    prefetch(&index[i]);  
  
for (i = 0; i<5; i++) {    /* Prolog 2 */  
    prefetch(&index[i+5]);  
    prefetch(&A[index[i]]);  
}  
  
for (i = 0; i<90; i++) { /* Steady State */  
    prefetch(&index[i+10]);  
    prefetch(&A[index[i+5]]);  
    sum += A[index[i]];  
}  
  
for (i = 90; i<95; i++) { /* Epilog 1 */  
    prefetch(&A[index[i+5]]);  
    sum += A[index[i]];  
}  
  
for (i = 95; i<100; i++) /* Epilog 2 */  
    sum += A[index[i]];
```

Summary of Results

Dense Matrix Code:

- eliminated 50% to 90% of memory stall time
- overheads remain low due to prefetching selectively
- significant improvements in overall performance (6 over 45%)

Indirections, Sparse Matrix Code:

- expanded coverage to handle some important cases

Prefetching for Arrays: Concluding Remarks

- Demonstrated that software prefetching is effective
 - selective prefetching to eliminate overhead
 - dense matrices and indirections / sparse matrices
 - uniprocessors and multiprocessors
- Hardware should focus on providing sufficient memory bandwidth

Prefetching for Recursive Data Structures

Recursive Data Structures

- Examples:
 - linked lists, trees, graphs, ...
- A common method of building large data structures
 - especially in non-numeric programs
- Cache miss behavior is a concern because:
 - large data set with respect to the cache size
 - temporal locality may be poor
 - little spatial locality among consecutively-accessed nodes

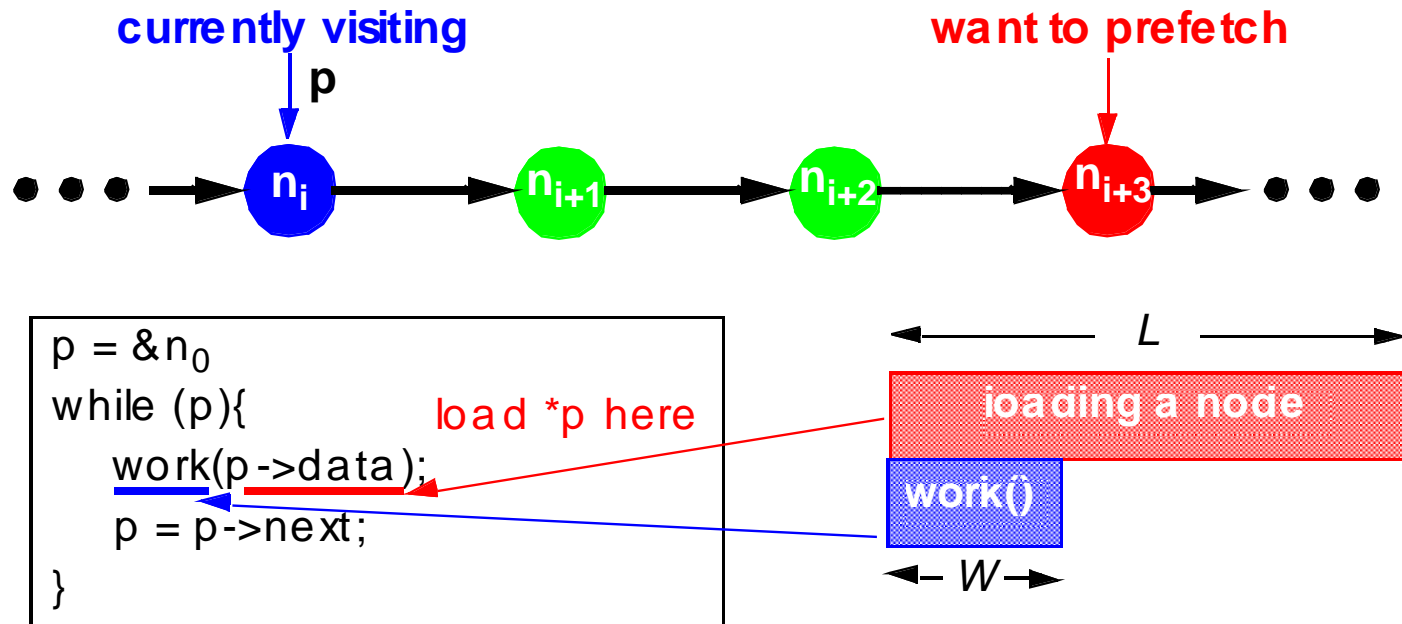
Goal:

- Automatic Compiler-Based Prefetching for Recursive Data Structures

Overview

- Challenges in Prefetching Recursive Data Structures
- Three Prefetching Algorithms
- Experimental Results
- Conclusions

Scheduling Prefetches for Recursive Data Structures

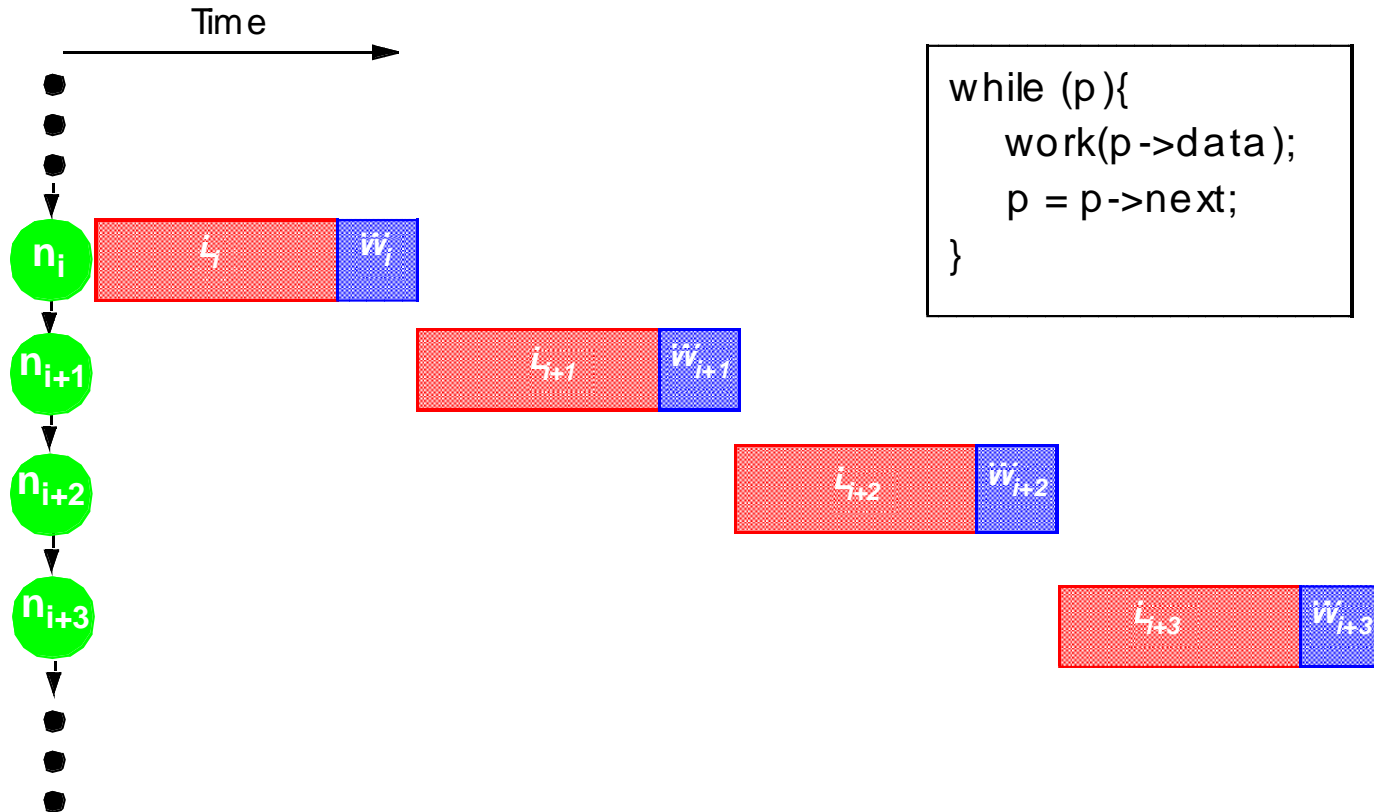


Our Goal: *fully hide latency*

– thus achieving fastest possible computation rate of $1/W$

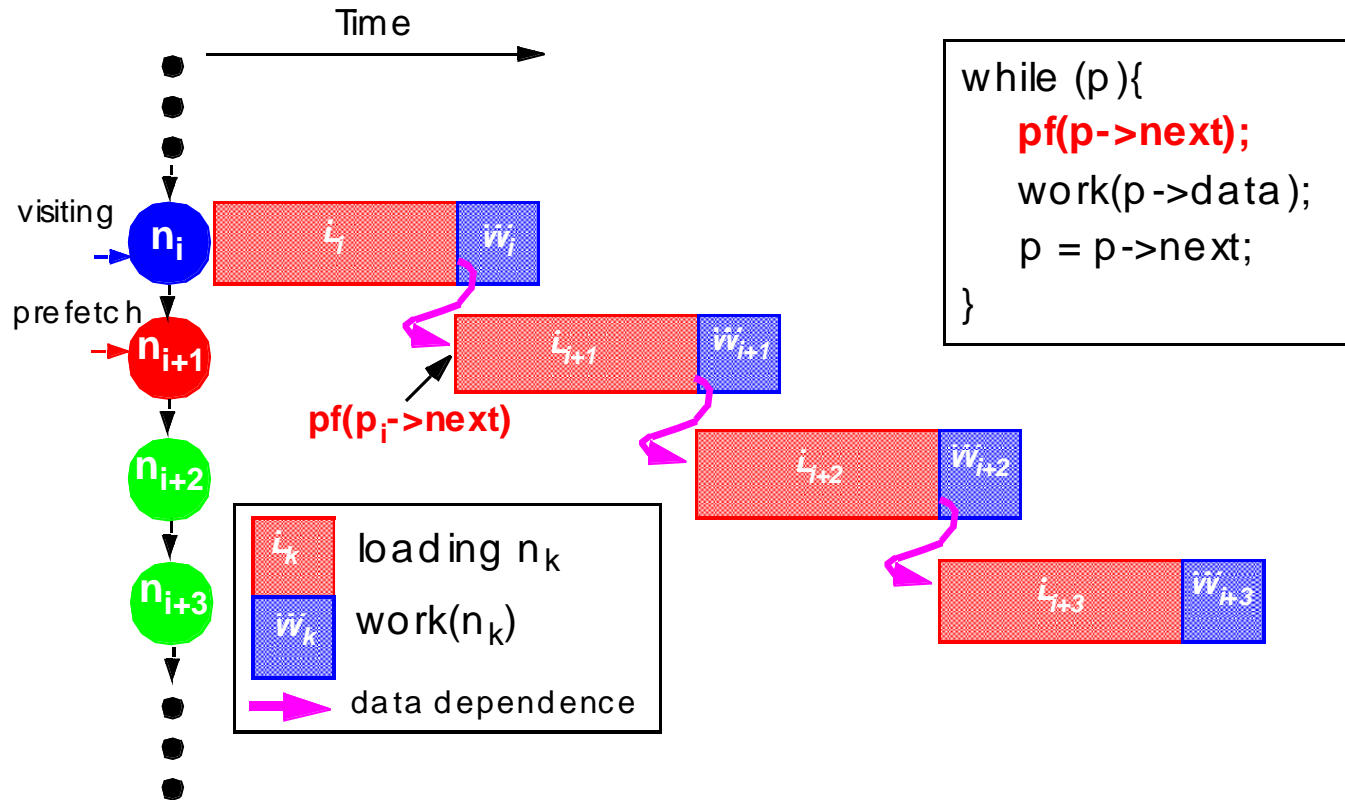
- e.g., if $L = 3W$, we must prefetch 3 nodes ahead to achieve this

Performance without Prefetching



$$\text{computation rate} = 1 / (L+W)$$

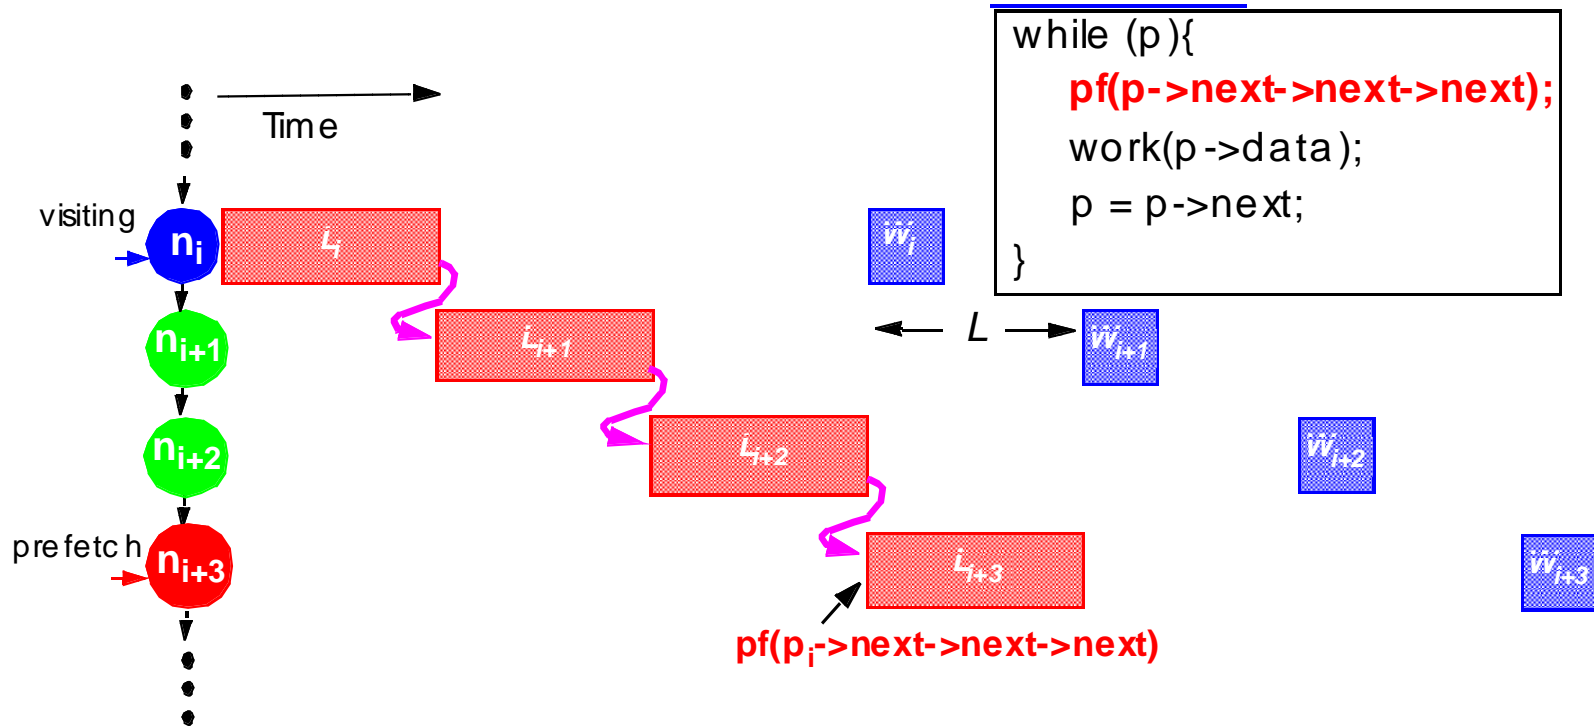
Prefetching One Node Ahead



- Computation is overlapped with memory accesses

computation rate = $1/L$

Prefetching Three Nodes Ahead




computation rate does not improve (still = 1/L)!

Pointer-Chasing Problem:

- any scheme which follows the pointer chain is limited to a rate of 1/L

Our Goal: Fully Hide Latency



```
while (p){
    pf(&n_{i+3});
    work(p->data);
    p = p->next;
}
```

- achieves the fastest possible computation rate of $1/W$

Overview

- Challenges in Prefetching Recursive Data Structures
- Three Prefetching Algorithms
 - Greedy Prefetching
 - History-Pointer Prefetching
 - Data-Linearization Prefetching
- Experimental Results
- Conclusions

Pointer-Chasing Problem

Key:

- n_i needs to know $\&n_{i+d}$ without referencing the $d-1$ intermediate nodes


Our proposals:

- use *existing* pointer(s) in n_i to approximate $\&n_{i+d}$
 - Greedy Prefetching
- add *new* pointer(s) to n_i to approximate $\&n_{i+d}$
 - History-Pointer Prefetching
- compute $\&n_{i+d}$ *directly* from $\&n_i$ (no ptr deref)
 - History-Pointer Prefetching

Greedy Prefetching

- Prefetch all neighboring nodes (simplified definition)
 - only one will be followed by the immediate control flow
 - hopefully, we will visit other neighbors later

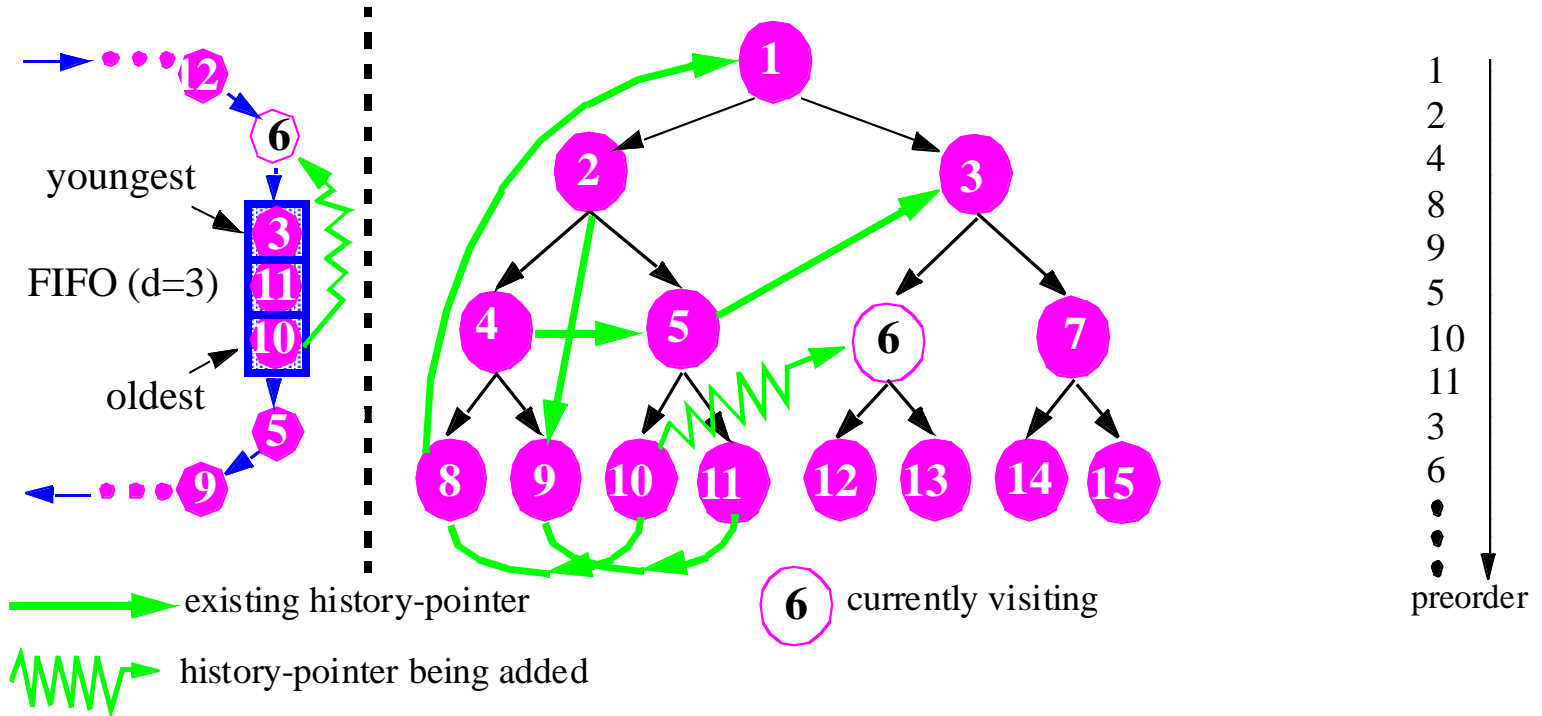
```
preorder(treeNode * t){  
    if (t != NULL){  
        pf(t->left);  
        pf(t->right);  
        process(t->data);  
        preorder(t->left);  
        preorder(t->right);  
    }  
}
```



- Reasonably effective in practice
- However, little control over the prefetching distance

History-Pointer Prefetching


- Add new pointer(s) to each node
 - history-pointers are obtained from some recent traversal



- Trade space & time for better control over prefetching distances

Data-Linearization Prefetching

- No pointer dereferences are required
- Map nodes close in the traversal to contiguous memory



Summary of Prefetching Algorithms

	<i>Greedy</i>	<i>History-Pointer</i>	<i>Data-Linearization</i>
<i>Control over Prefetching Distance</i>	little	more precise	more precise
<i>Applicability to Recursive Data Structures</i>	any RDS	revisited; changes only slowly	must have a major traversal order; changes only slowly
<i>Overhead in Preparing Prefetch Addresses</i>	none	space + time	none in practice
<i>Ease of Implementation</i>	relatively straightforward	more difficult	more difficulty

Conclusions

- Propose 3 schemes to overcome the pointer-chasing problem:
 - Greedy Prefetching
 - History-Pointer Prefetching
 - Data-Linearization Prefetching
- Automated greedy prefetching in SUIF
 - improves performance significantly for half of Olden
 - memory feedback can further reduce prefetch overhead
- The other 2 schemes can outperform greedy in some situations

CSC D70: Compiler Optimization Prefetching

Prof. Gennady Pekhimenko

University of Toronto

Winter 2021

*The content of this lecture is adapted from the lectures of
Todd Mowry and Phillip Gibbons*